

What the order book does not show you

Week 1. Post-only rejection, measured on 32 addresses over 16,244 polls.

Every figure here was measured on 20 September 2026. Where a number is not measured, it is not shown.

THE PROBLEM

A refused order leaves no trace

A maker sends a post-only order. It is meant to rest on the book and wait.
If it would cross the spread, the matching engine refuses it instead.

That refusal is invisible. No fill. No volume. No depth. The order never existed anywhere a person looks, and the book it was meant to join is thinner than the screen says.

What is visible

Fills. Volume. Depth.

What is not

Every order the engine turned away before it rested.

The same number, two opposite meanings

A trader reading the book

The depth shown is not the depth available. Orders that were refused never arrived, so the quote is built on liquidity that was turned away.

A high rate means the book is thinner than it appears.

A protocol routing a user's order

You send someone else's order through a builder code. A refused post-only did not rest and did not fill. The user got nothing.

A high rate means your users are being dropped.

For the maker the refusal is the mechanism working, saving it a taker fee. For the person whose order it was, it is a failure. That distinction is the reading this project adds.

WHERE IT SHOWS UP

Measured, not asserted

0.17%

the median maker in the median market

across 40 markets carrying 50,000 post-only orders or more

100%

one address, one market, 252,994 orders

refused every time in xyz:SP500, while its rate across all markets reads 27.23%

88x

BTC pooled against BTC median

58.4% pooled, 0.66% for the median maker, because 2 of 23 makers place most of the orders

WHAT IT LOOKS LIKE IN PLACE

On the page you were already reading



Post-only rejection

0x7fda...17d1 · measured by Tnega

Tnega

Who is behind this address is not established


The venue reports an ordinary account with no approved builder. That is what a single trader looks like and also what a platform holding many customers in one account looks like, and nothing public separates the two. Read the rate as the behaviour of an account, not of a person.

1.4%

QUOTING

Almost every post-only order rests on the book. This is what market making looks like.

LAST 49 HOURS



POST-ONLY REJECTED

11,067 of 784,648

POLLS STORED

396

NEWEST ORDER SEEN

6 min ago

A rejected post-only order never rests on the book, so it provides no liquidity and leaves no trace in fills or volume.

This does not say whether the address is any good or making money. A low rate means its quotes reach the book, which is what quoting looks like rather than a score. The tracked set holds addresses that quote cleanly and lose, and addresses refused constantly that gain.

ON HYPERCORE, READ ON CHAIN

Short 578.3531 BTC

entry \$79,811.82 · mark \$81,497

-\$974,626

Short 22,043.2243 ETH

entry \$2,399.48 · mark \$2,641.56

-\$5,336,326

Short 53,644.36 SOL

entry \$102.38 · mark \$111.75

-\$502,622

Short 2,447,671 XRP

entry \$1.3397 · mark \$1.4525

-\$276,179

Short 162.98512 MYDE

\$2,030,536

The rate, the hours behind it, and what the address holds on HyperCore. Where a rate cannot be stated, the panel gives the reason instead of a number.

Seventy times apart

Pooled across every order on the venue, the rejection rate is **12.5%**:
2,317,205 of 18,469,277 post-only orders.

The median maker in the median market sees **0.17%**.

Both are correct. They answer different questions. A protocol that reads the pooled figure and concludes its users will be refused one time in eight is wrong by a factor of roughly seventy, because its flow behaves like the median maker and not like the two addresses that dominate the pool.

Pooled

Every order on the book, counted together. Dominated by whoever submits most.

Median maker

One value per address, then the middle one. What a participant like you actually sees.

Four pieces

Collectors

A REST poller every 15 minutes over the tracked set, and a WebSocket worker bucketing order updates at ten seconds. 16,244 polls so far.

A reader contract

Deployed on HyperEVM, reading HyperCore through the precompiles, so positions and marks come off chain rather than from an API that could be wrong.

The panel

One rate, the hours behind it, and the counts. Where the observations are too few or the data has aged, the reason replaces the number.

The extension

The same reading injected on nine sites, so it reaches a reader on the page they were already on rather than requiring a visit here. Seven are block explorers; the other two are Hyperliquid's own app and the registry indexer.

My own numbers were measuring the wrong thing

A count inflated by a fifth

The venue returns a rolling window of recent orders and ignores the date range asked for. When an address is quiet, consecutive polls return the same records and each poll stored them again.

About 21% of the 35.6 million stored records are the same orders counted twice, leaving roughly 28 million distinct. The rates were unaffected, because numerator and denominator inflate together. The total was not.

A profit finding withdrawn after publication

I reported that the venue's leaderboard contradicted itself, on the grounds that month profit exceeded all-time profit on 51.3% of rows.

Profit is signed. An account down since it opened and up this month satisfies that comparison legitimately, and 87.9% of the rows I flagged had a negative all-time figure. The test detected a minus sign.

Both are written up in the repository rather than quietly corrected, because a measurement project that hides its own corrections is asking for trust it has not earned.